

Exercise 4 - LDA

Munther Odeh

Timo Marks

Overall very nice! However, in the future, please also hand in the ipython notebook file so I can test your code.

Two small things I noticed, or maybe overlooked since reading a notebook in pdf form is not that great, you did not normalize your data, and you did not add any bias weights. LDA assumes that the data underlies a normal distribution, which it is not. On the other hand, LDA is actually quite robust, especially for simple binary classification like our task, so it will only increase the performance by another percent or so. Adding a bias is also not that important for LDA (but it would be for Least Squares!), because the data should be normalized beforehand.

Total 10/10 points. Please upload the notebook when you read this.